

# EnviroHealth-Monitor: Technical Project Report

## A Distributed Real-Time Pipeline for Urban Environmental Analytics

Date: January 2026

Institution: MSc Computer Science & Data Science @ ESILV

Project Demo: <https://youtu.be/IVGqxV3CXvA>

### The Team

The following team members contributed to the implementation and demonstration of this distributed system:

- Hussnain Amanat Ali: Project Lead & Spark Logic
  - Muhammad Irfan: Infrastructure & YARN Management
  - Ayman Berri: Monitoring & Dashboard Visualization
- 

## 1. Executive Summary

EnviroHealth-Monitor addresses the demand for real-time monitoring of urban environmental conditions in metropolitan areas like Paris. By moving away from traditional high-latency batch processing, this project delivers an end-to-end streaming architecture with sub-5-second latency. The system provides a 360-degree view of both Infrastructure Health and Environmental Analytics.

---

## 2. Distributed Infrastructure

Unlike standard local deployments, this project utilizes a professional cloud-based distributed architecture:

- Cluster Scale: A 6-node Hadoop cluster (1 Master, 5 Workers) deployed on Microsoft Azure.
  - Hardware Capacity: Total cluster power of 48 GB RAM and 48 vCores (8GB/8-vCores per node).
  - Resource Orchestration: Managed via Apache YARN, ensuring all nodes maintain a "RUNNING" state during high-compute Spark tasks.
- 

## 3. The Data Pipeline

The system is designed as a decoupled, four-stage distributed pipeline to ensure high availability and fault tolerance.

#### A. Ingestion Layer

Real-time data is ingested from the OpenAQ and OpenWeather APIs. Python-based producers simulate IoT sensors by publishing raw JSON messages to dedicated Kafka topics.

#### B. Processing Layer (Apache Spark)

Apache Spark Structured Streaming serves as the core processing engine. It performs:

- Strict Schema Enforcement: Transforming raw JSON into typed DataFrames.
- Windowed Aggregations: Calculating real-time pollution averages (PM2.5) every few seconds.
- Correlation: Joining weather and air quality streams to identify health risk patterns.

#### C. Storage & Data Lake

- Time-Series Storage: Processed metrics are written to InfluxDB for temporal analysis.
  - Cold Storage: Raw and processed data is permanently saved in HDFS, partitioned by year and month for future batch auditing.
- 

### 4. Monitoring & Visualization (TIG Stack)

The project implements the TIG Stack (Telegraf, InfluxDB, Grafana) to monitor system and application performance simultaneously:

- Infrastructure Monitoring: Telegraf agents collect real-time RAM usage from the Azure VMs, visualized in a "Usage Graph" to prevent cluster bottlenecks.
  - Functional Visualization: Grafana dashboards display live "Pollution Deep-Dives" and "Worker Audits" (Batch Count: 8.11), proving pipeline integrity.
- 

### 5. Technical Challenges Solved

1. OS Integration: Successfully deployed the full stack on Ubuntu 24.04 LTS, overcoming modern repository and GPG key compatibility issues.

2. **Distributed Consistency:** Synchronized configuration files (`yarn-site.xml`, `core-site.xml`) across 6 nodes to ensure seamless task distribution.
  3. **Real-time Synchronization:** Coordinated separate API streams into a unified Kafka-Spark flow for multi-dimensional analysis.
- 

## 6. Conclusion

**EnviroHealth-Monitor** successfully demonstrates a modern distributed Big Data stack. The system achieves low latency and high reliability, proving that cloud-distributed clusters can provide vital real-time environmental insights while maintaining robust infrastructure monitoring.